

Homework 2

Yiming Gao (NetID: yimingg2)

2017/2/5

Contents

Exercise 1	1
Exercise 2	4
Exercise 3	5

Exercise 1

In this exercise, we investigate the bias-variance tradeoff when estimating the function f defined below.

```
f = function(x1, x2) {  
  x1 ^ 3 + x2 ^ 3  
}
```

The following code defines the data generating process and should we used to simulate data.

```
get_sim_data = function(f, sample_size = 100) {  
  x1 = runif(n = sample_size, min = -1, max = 1)  
  x2 = runif(n = sample_size, min = -1, max = 1)  
  y = f(x1, x2) + rnorm(n = sample_size, mean = 0, sd = 0.5)  
  data.frame(x1, x2, y)  
}
```

Use simulation to investigate the bias and variance of *five* models at the point $\mathbf{x} = (x_1, x_2) = (0.75, 0.95)$. The five models are of the form

- $y \sim \text{poly}(x_1, \text{degree} = k) + \text{poly}(x_2, \text{degree} = k)$

for $k = 1, 2, 3, 4, 5$. Use 500 simulated samples each of size 200. Before performing the simulations, I set the seed as 650379994.

```
uin = 650379994  
set.seed(uin)
```

```

n_sims = 500
x01 = 0.75
x02 = 0.95
pred = matrix(0, nrow = n_sims, ncol = 5)

for (i in 1:n_sims){
  sim_data = get_sim_data(f, sample_size = 200)

  fit1 = lm(y ~ poly(x1, degree = 1) + poly(x2, degree = 1), data = sim_data)
  fit2 = lm(y ~ poly(x1, degree = 2) + poly(x2, degree = 2), data = sim_data)
  fit3 = lm(y ~ poly(x1, degree = 3) + poly(x2, degree = 3), data = sim_data)
  fit4 = lm(y ~ poly(x1, degree = 4) + poly(x2, degree = 4), data = sim_data)
  fit5 = lm(y ~ poly(x1, degree = 5) + poly(x2, degree = 5), data = sim_data)

  pred[i, ] <- c(
    predict(fit1, newdata = data.frame(x1 = x01, x2 = x02)),
    predict(fit2, newdata = data.frame(x1 = x01, x2 = x02)),
    predict(fit3, newdata = data.frame(x1 = x01, x2 = x02)),
    predict(fit4, newdata = data.frame(x1 = x01, x2 = x02)),
    predict(fit5, newdata = data.frame(x1 = x01, x2 = x02))
  )
}

```

We simulate 500 samples each of size 200 for building our models.

```

# Tradeoff
eps = rnorm(n = n_sims, mean = 0, sd = 0.3)
y0 = f(x01, x02) + eps

# get bias
get_bias = function(estimate, truth) {
  mean(estimate - truth)
}

# get mean square error
get_mse = function(estimate, truth) {
  mean((estimate - truth) ^ 2)
}

# suppress scientific notation

```


```
options(scipen=999)
bias = apply(pred, 2, get_bias, y0)
variance = apply(pred, 2, var)
mse = apply(pred, 2, get_mse, y0)
```

Model	Squared Bias	Variance	MSE
fit1	0.0800646	0.0084093	0.173899
fit2	0.0832004	0.0171158	0.1868814
fit3	0.0007602	0.0189415	0.1034944
fit4	0.000826	0.0255703	0.1114743
fit5	0.0007265	0.0327806	0.1168097

We summarize the table above, which contains **Model**, **squared bias**, **variance** and **MSE** of the estimates. We notice that fit3 and fit5 has the lowest squared bias. However, fit5 has the highest variance among five models.

```
mydata = data.frame("degree" = c(1:5), "Bias" = bias^2, "Variance" = variance)

library(ggplot2)
ggplot(mydata, aes(degree))+
  geom_smooth(aes(y = Bias, colour = "Squared bias"))+
  geom_smooth(aes(y = Variance, colour = "Variance")) +
  xlab("degree") + ylab(label = "Value") +
  # add legend
  scale_colour_manual(values = c("hotpink", "dodgerblue"))
```



We plot **Squared bias** and **Variance** values against the **degree** of the polynomials used. From the plot, we know that squared bias decreases significantly when degree increases to 3. The variance values appear to have a positive linear trend as degree increasing. So we may consider **model 3** achieves a balance between bias and variance.

Exercise 2

We need to find a model that satisfies:

- Train RMSE less than 1.08
- Test RMSE less than 1.01

The codes for building model and some relevant functions are as follows:

```
# user-defined function to calculate rmse
rmse = function(actual, predicted) {
  sqrt(mean((actual - predicted) ^ 2))
}

get_rmse = function(model, data, response) {
```

```

rmse(actual = data[, response],
      predicted = predict(model, data))
}

model = lm(y ~ . + x1:x2 + I(x3^2) + I(x4^3), data = train_data)

complexity = length(coef(model)) - 1
train_rmse = get_rmse(model, data = train_data, response = "y")
test_rmse = get_rmse(model, data = test_data, response = "y")

```

After a number of trials, we get our final model and its metrics:

- $y \sim . + x1:x2 + I(x3^2) + I(x4^3)$

Model	Train RMSE	Test RMSE	Number of parameters
model	1.03	0.978	7

Train RMSE is $1.030 < 1.08$. Test RMSE is $0.978 < 1.01$.

Exercise 3

We want to build a logistic regression model using `mpg` as the response.

We use the training data to train a classifier which achieves:

- Train Accuracy greater than 0.89
- Test Accuracy greater than 0.89

```

library(caret)
model_glm = glm(mpg ~ ., data = auto_train_data, family = "binomial")

# obtain the predicted probabilities, use 0.5 as cutoff
glm_train_pred = ifelse(predict.glm(model_glm, newdata = auto_train_data,
                                   type = "response") > 0.5, 1, 0)
glm_test_pred = ifelse(predict.glm(model_glm, newdata = auto_test_data,
                                   type = "response") > 0.5, 1, 0)

train_accuracy = mean(glm_train_pred == auto_train_data$mpg) # 0.897
test_accuracy = mean(glm_test_pred == auto_test_data$mpg) # 0.924

```

```
# confusion matrix for test data
test_tab = table(predicted = glm_test_pred, actual = auto_test_data$mpg)

# sensitivity and specificity
test_mat = confusionMatrix(test_tab, positive = "1")
```

We finally train the following classifier using training set:

- $\text{mpg} \sim \text{cylinders} + \text{displacement} + \text{horsepower} + \text{weight} + \text{acceleration} + \text{year}$

It has

- Train Accuracy = $0.897 > 0.89$
- Test Accuracy = $0.924 > 0.89$

The confusion matrix for test data is:

```
##          actual
## predicted  0  1
##          0 44  1
##          1  6 41
```

And other metrics for the test data are summarized in the following table:

Test sensitivity	Test specificity
0.976	0.88